Credit Risk Management and Portfolio Optimization:

1. Abstract :

This report details the development and implementation of an end-to-end credit risk management framework utilizing machine learning, specifically addressing the Home Credit Default Risk challenge. The project aimed to enhance traditional credit assessment methods by predicting loan default risk, identifying high-risk customers, and optimizing lending portfolios. A LightGBM model, achieving an ROC AUC of approximately 0.777, forms the core of the predictive engine. Beyond mere prediction, the framework integrates advanced components such as a composite Financial Stress Score, data-driven customer personas (e.g., Prime Borrower, Financially Stressed Customer), an Early Warning System, and K-Means-based customer segmentation. Crucially, SHAP (SHapley Additive exPlanations) analysis was employed to ensure model transparency and explainability, a critical aspect for regulatory compliance and business trust. The solution demonstrated significant business impact by reducing overall portfolio default risk from 8.07% to 4.94% through optimized lending strategies and risk-based pricing. This comprehensive approach transforms raw transactional data into actionable insights, enabling more informed and sustainable lending decisions.

2. Introduction :

Financial institutions globally face substantial financial losses due to loan defaults. Traditional credit scoring models often struggle to capture the intricate, non-linear relationships between diverse customer attributes, credit history, and repayment behaviors. This project was initiated to address these limitations by developing a robust, data-driven framework for credit risk management. The primary objectives included accurate loan default prediction, proactive identification of risky customers, development of interpretable risk indicators, and strategic portfolio optimization. The Home Credit Default Risk dataset served as the foundation for this initiative, providing a rich, multi-source data environment to build and validate the proposed machine learning solution.

3. Data and Feature Engineering :

3.1 Dataset Description and Integration

The project leveraged the Home Credit Default Risk dataset, which comprises multiple interconnected tables including application , bureau , previous_application , installments_payments , POS_CASH_balance , and credit_card_balance . These disparate sources were integrated and aggregated using unique customer identifiers to construct a master dataset. The final engineered dataset encompassed approximately 307,510 customer records and 183 features, with the target variable indicating loan default status.

3.2 Data Preprocessing

A rigorous data preprocessing pipeline was implemented to ensure data quality and model readiness. This involved: **Missing Value Imputation**: Strategic handling of missing values using domain-appropriate techniques (e.g., mean, median, mode imputation, or advanced methods like MICE).

Outlier Treatment: Identification and mitigation of outliers to prevent undue influence on model training.

Numerical Variable Cleaning: Standardization or normalization of numerical features to ensure consistent scaling.

Categorical Feature Encoding: Transformation of categorical variables into numerical representations suitable for machine learning algorithms (e.g., One-Hot Encoding, Label Encoding).

3.3 Feature Engineering

Extensive feature engineering was performed to extract meaningful, business-relevant signals from the raw data. This process transformed transactional and demographic data into predictive indicators. Key engineered features included:

Age and Employment Years: Derived features capturing the stability and experience of applicants.

Loan and Annuity Burden: Ratios indicating the proportion of income allocated to loan repayments.

Credit Dependency Ratio: Metric reflecting reliance on credit.

Goods-to-Credit Ratio: Insight into the value of purchased goods relative to credit extended.

Debt Burden Indicators: Comprehensive metrics assessing overall debt load.

Overdue Ratio: Indicators of past payment delinquency.

Income per Family Member: A normalized income metric accounting for household size.
Financial Stress Score: A composite score quantifying repayment pressure, incorporating income, debt burden, loan obligations,
and repayment characteristics. This score enabled categorization of customers into distinct stress bands for enhanced risk
monitoring.
Customer Value Score & Risk Score: Proprietary scores designed to capture customer long-term value and inherent risk profiles.
These engineered features were crucial in translating raw data into actionable risk indicators, significantly enhancing the model’s
predictive power and interpretability.

4. Methodology

4.1 Machine Learning Model: LightGBM

LightGBM (Light Gradient Boosting Machine) was selected as the primary predictive model due to its proven effectiveness with structured tabular data, efficiency in handling large datasets, robust performance, and native support for missing values. Its gradient boosting framework constructs an ensemble of decision trees, iteratively improving predictions by correcting errors from previous trees. The model was trained to predict the probability of loan default, achieving an ROC-AUC of approximately 0.777.

4.2 Customer Personas

To bridge the gap between technical model outputs and business understanding, a customer persona framework was developed. This
involved classifying customers into distinct, interpretable categories based on their financial behavior and risk profiles. Examples of
identified personas include:
Prime Borrower: Low-risk, financially stable customers.
Growth Borrower: Customers with potential for growth, requiring careful management.
Standard Customer: Average risk profile.
Financially Stressed Customer: High-risk customers exhibiting signs of repayment difficulty.
These personas enable lending teams to tailor strategies and communications more effectively.

4.3 Early Warning System (EWS)

An Early Warning System was implemented to facilitate proactive risk management. Customers were dynamically categorized into three groups based on their evolving risk indicators:

Green: Low-risk borrowers requiring minimal monitoring.

Amber: Customers showing early signs of elevated risk, warranting closer observation.

Red: High-risk customers requiring immediate intervention to mitigate potential default.

Empirical analysis confirmed a significant increase in observed default rates from Green to Red categories, validating the EWS' s efficacy in identifying deteriorating risk profiles.

4.4 Customer Segmentation

Unsupervised learning, specifically K-Means clustering, was applied to segment customers into groups with distinct financial characteristics. This segmentation provided a deeper understanding of the customer base beyond individual risk scores. Identified segments included:

Prime High Value: Customers with excellent credit and high transaction volumes.

Mature Stable: Established customers with consistent financial behavior.

Credit Dependent: Customers with a higher reliance on credit.

Young High Risk: Newer customers with potentially volatile financial profiles.

These segments enable differentiated lending strategies, product offerings, and targeted marketing campaigns.

5. *Model Performance and Explainability*:

5.1 Model Performance

The LightGBM model achieved a Receiver Operating Characteristic Area Under the Curve (ROC-AUC) of approximately 0.777. This metric indicates a strong ability to distinguish between defaulting and non-defaulting customers, outperforming many traditional credit scoring methods. The model' s performance was rigorously validated using cross-validation techniques and held-out test sets to ensure generalization capability.

5.2 Model Explainability using SHAP

In the context of financial services, model interpretability is paramount for regulatory compliance, auditability, and building trust with stakeholders. SHAP (SHapley Additive exPlanations) values were utilized to explain individual model predictions and global feature importance. SHAP analysis revealed that external risk scores, age-related features, employment characteristics, and credit history metrics were the most significant contributors to the model’ s predictions. This explainability enhances transparency, allowing business users to understand *why* a particular lending decision was made, rather than just *what* the decision was.

6. Business Impact and Optimization :

6.1 Solution Alignment with Problem Statement

The developed framework directly addresses the core problem of inefficient credit risk assessment and suboptimal lending decisions.
By integrating predictive modeling with business analytics, the solution:

1. **Predicts Default Probability**: Provides accurate, data-driven forecasts of loan default.
2. **Identifies High-Risk Customers**: Enables proactive intervention before defaults occur.
3. **Provides Explainable Predictions**: Fosters trust and facilitates regulatory compliance.
4. **Segments Customers**: Creates actionable business groups for tailored strategies.
5. **Supports Risk-Based Pricing**: Aligns interest rates with customer risk profiles.
6. **Reduces Portfolio Risk**: Optimizes overall lending exposure.
7. **Enables Data-Driven Strategies**: Shifts from rule-based to intelligent, adaptive lending.

6.2 Portfolio Optimization

Beyond individual default prediction, a key focus of this project was portfolio-level risk optimization. By leveraging customer segments and risk scores, various lending scenarios were simulated. The optimized portfolio strategy resulted in a significant reduction of overall portfolio default risk from approximately 8.07% to 4.94%. This represents a substantial decrease in expected risk exposure, directly contributing to improved financial stability and profitability for the lending institution.

6.3 Risk-Based Pricing Strategy

To further align pricing with risk, a dynamic risk-based pricing strategy was recommended, leveraging the identified customer segments:

Customer Segment	Recommended Interest Rate
Prime High Value	10%
Mature Stable	11%
Credit Dependent	14%
Young High Risk	18%

This strategy ensures that interest rates are commensurate with the perceived risk of each customer segment, promoting sustainable profitability while managing risk effectively.

7. Deployment Considerations and Future Work

As a CSE Student, deployment and ongoing maintenance are critical considerations. For this framework, future enhancements would focus on operationalizing the models and expanding their capabilities:

Real-Time Scoring Systems: Transitioning from batch processing to real-time inference for immediate credit decisions, requiring robust API development and low-latency infrastructure.

Web-based Deployment: Developing a user-friendly interface (e.g., using Streamlit or a custom web application) for business users to interact with the model, visualize predictions, and simulate scenarios.

Integration with Banking Workflows: Seamless integration of the ML framework into existing core banking systems and decision engines, necessitating careful API design and data synchronization strategies.

Deep Learning Approaches: Exploring advanced deep learning architectures (e.g., Recurrent Neural Networks for sequential data, or Transformers for complex feature interactions) to potentially capture more nuanced patterns in credit behavior, especially with larger and more diverse datasets [4].

Continuous Model Monitoring and Retraining: Implementing automated pipelines for monitoring model performance drift, data quality, and concept drift. Regular retraining with fresh data is essential to maintain model accuracy and relevance over time [7].

Automated Risk Recommendation Systems: Developing systems that not only predict risk but also provide automated, context aware recommendations for risk mitigation strategies or personalized product offerings.

Addressing challenges such as data privacy, ethical considerations, and regulatory compliance will be paramount during deployment, particularly when dealing with sensitive financial data [5]. Robust MLOps practices, including version control for models and data, automated testing, and continuous integration/continuous deployment (CI/CD) pipelines, will be essential for successful and sustainable operation.

8. Conclusion

This project successfully demonstrates the transformative potential of machine learning in credit risk management. By evolving from a basic default prediction task to a comprehensive framework encompassing predictive modeling, explainability, segmentation, risk scoring, and portfolio optimization, the solution provides actionable insights for financial institutions. The integration of advanced ML techniques with business-centric interpretations empowers lenders to make smarter, data-driven decisions, ultimately leading to reduced risk exposure and enhanced profitability. The framework serves as a robust foundation for future innovations in intelligent credit risk assessment.

9. References

[1] Kaggle. *Home Credit Default Risk*. Available at: https://www.kaggle.com/c/home-credit-default-risk [2] Business Science. *How To Complete A Kaggle Competition In 30 Minutes*. (2018). Available at: https://www.business-science.io/business/2018/08/07/kaggle competition-home-credit-default-risk.html [3] Medium. *HOME CREDIT DEFAULT RISK — An End to End ML Case Study*. (2020). Available at: https://medium.com/thecyphy/home-credit-default-risk-an-end-to-end-ml-case-study-part-1-introduction-and-eda-3bfe3c7ddd7a [4] Abi, R. (2025). *Machine learning for credit scoring and loan default prediction using behavioral and transactional financial data*. World Journal of Advanced Research and Reviews. Available at: https://www.researchgate.net/profile/Roland Abi/publication/393147041_Machine_learning_for_credit_scoring_and_loan_default_prediction_using_behavioral_and_transactional_financia

learning-for-credit-scoring-and-loan-default-prediction-using-behavioral-and-transactional-financial-data.pdf [5] Faheem, M. A. (2021). *AI-driven risk assessment models: Revolutionizing credit scoring and default prediction*. Iconic Research And Engineering Journals.
Available at: https://www.researchgate.net/profile/Muhammad-Ashraf-Faheem/publication/385074551_AI Driven_Risk_Assessment_Models_Revolutionizing_Credit_Scoring_and_Default_Prediction/links/6713e02e035917754c09413e/AI Driven-Risk-Assessment-Models-Revolutionizing-Credit-Scoring-and-Default-Prediction.pdf [6] Bhatore, S., Mohan, L., & Reddy, Y. R. (2020). *Machine learning techniques for credit risk evaluation: a systematic literature review*. International Journal of Banking and Financial Technology. Available at: https://link.springer.com/article/10.1007/s42786-020-00020-3 [7] IEEE Xplore. *Software engineering for machine learning: A case study*. (2019). Available at: https://ieeexplore.ieee.org/abstract/document/8804457/

Team Member Name :

1. Kaustav Maji
2. Anushkan Panda